



Q4_0 F16

(a) Prefill



(b) Decode



Qwen2.5-1.5B, Galaxy S25 (6 threads). * NPU+F16 falls back to CPU companion (NPU dispatch is limited to Q4_0 weight matmul).