




Qwen2.5-1.5B, Galaxy S25 (6 threads). All cells dispatch on the named backend; F16 weights are mirrored to rpcmem for NPU dispatch. F16-on-NPU is memory-bound (~3.3 GB rpcmem) and hence slowest.