




Qwen2.5-1.5B decode shapes, Galaxy S25. HTP latency is floored at ~100 us (FastRPC round-trip + DMA-BUF coherency) regardless of compute; only ops whose compute exceeds the floor (softmax, flash\_attn, matmul) benefit from NPU offload.